

UNIVERSITÉ LIBRE DE BRUXELLES

DATABASE SYSTEMS ARCHITECTURE

PROJECT ASSIGNMENT

Algorithms in Secondary Memory

Author:

Pierre-Alexandre BOURDAIS

Matricule ULB: 000458467

Shafagh KASHEF

Matricule ULB: 000455644

Keneth UBEDA

Matricule ULB: 000453317

Group: 12

Supervisor:

Prof. S.VANSUMMEREN

January 7, 2018



Contents

1	Introduction and Environment	2
1.1	Abstract	2
2	Observations on streams	3
2.1	First I/O mechanism	3
2.1.1	Implementation	3
2.1.2	Expected behavior	3
2.2	Second I/O mechanism	3
2.2.1	Implementation	3
2.2.2	Expected behavior	4
2.3	Third I/O mechanism	4
2.3.1	Implementation	4
2.3.2	Expected behavior	4
2.4	Fourth I/O mechanism	5
2.4.1	Implementation	5
2.4.2	Expected behavior	5
2.5	Experimental observations	5
2.6	Configuration	5
2.6.1	Write-Benchmark	6
2.6.2	Read-Benchmark	16
2.7	Discussion of expected behavior vs experimental observations . .	21
3	Observations on multi-way merge sort	21
3.1	Implementation	21
3.2	Expected behavior	23
3.3	Experimental observations	23
3.4	Discussion of expected behavior vs experimental observations . .	25
4	Overall conclusion	25
5	References	27

1 Introduction and Environment

1.1 Abstract

The goal of the project is to implement an external-memory merge-sort algorithm and examine its performance under different parameters (data volume, size of the buffer,...). Different read data and write data functions using such as memory mapping and buffer are implemented in this project and these implementations helped to implement the external-memory merge-sort algorithm. The language used to implement the algorithms is Java and the IDE is NetBeans. In this Project very large relations with Two-Phase, Multiway Merge-Sort (TPMMS) algorithm is sorted.

The environment of the project:

Machine type: MacBook Air
Hard Disk Type: Macintosh HD
Operating System: MacOS
Memory : 8 Go 1600 MHz DDR3
Total Memory Available: 10Go
Programming Language: Java 8
Libraries : JMH

The merge sort implementation is able to sort disk files consisting of 32-bit integers. And reads data from, and writes data to disk. Therefore stream classes are developed that can sequentially read and write a file consisting of 32-bit Integers. Two types of streams are used: input streams and output streams.

The input stream supports the following operation:

- open : open an existing file for reading
- read next : read the next element from the stream
- end of stream : a Boolean operation that returns true if the end of stream has been reached

The output stream supports the following operations:

- create : create a new file
- write : write an element to the stream
- close : close the stream

2 Observations on streams

2.1 First I/O mechanism

2.1.1 Implementation

The first I/O mechanism is quite simple as it requires to read and write one element at a time using the read and write system calls.

- Read function mimic read system call by calling `readInt()` on a `java.io.DataInputStream` that is wrapped directly around a `java.io.FileInputStream`.
- Write function mimic write system call by calling `write` on a `java.io.DataOutputStream` that is wrapped directly around a `java.io.FileOutputStream`.

2.1.2 Expected behavior

Before running any test, a cost formula needs to be build for this implementation. The cost estimates the total number of I/O's that need to be done in function of N and k .

- k - the number of streams, where each stream is a file
- N - the data volumes
- $B(R)$ - the number of blocks that R occupies on disk

The algorithm needs to read $B(R)$ blocks, write $B(R)$ blocks and open $B(R)$ elements.

$$\begin{aligned}\text{Cost(1st I/O)} &: N.B(R) + N.B(R) + k.B(R) \\ \text{Cost(1st I/O)} &: 2.N.B(R) + k.B(R) \\ \text{Cost(1st I/O)} &: (2.N + k).B(R)\end{aligned}$$

2.2 Second I/O mechanism

2.2.1 Implementation

The second I/O mechanism need an other implementation of read and write functions. This time the read function has to read element by element each one with a size of 32-bit from a stream and stores them in a buffer. And the write function has to write one element (32-bit) from the buffer. More over this mechanism need also a new open function in order to use a buffer.

- Read function mimic fread system call by calling readInt() on a java.io.BufferedReader that itself is wrapped around a java.io.FileInputStream.
- Write function mimic fwrite system call by calling writeInt on a java.io.BufferedOutputStream that is wrapped directly around a java.io.FileOutputStream.

2.2.2 Expected behavior

Before running any test, a cost formula needs to be build for this implementation. The cost estimates the total number of I/O's that need to be done in function of N and k.

- k - the number of streams, where each stream is a file
- N - the data volumes
- N(R) - the number of blocks that R occupies on disk

The algorithm needs to read B(R) blocks, write B(R) blocks from a buffer and open B(R) elements.

$$\begin{aligned} \text{Cost(2nd I/O)} &: N.B(R) + k.B(R) \\ \text{Cost(2nd I/O)} &: (N+k)B(R) \end{aligned}$$

2.3 Third I/O mechanism

2.3.1 Implementation

The third I/O mechanism is almost the same as the first one implemented. Indeed in this implementation of read and write function it's just need to add a buffer of size B in internal memory. Thus once the buffer becomes empty/full the next B elements are read/written from/to the file.

2.3.2 Expected behavior

Before running any test, a cost formula needs to be build for this implementation. The cost estimates the total number of I/O's that need to be done in function of N, B and k.

- k - the number of streams, where each stream is a file
- N - the size of the input file, measured in number of 32-bit integers.
- B(R) - the number of blocks that R occupies on disk
- b is the size of the buffer in internal memory

$$\text{Cost(3rd I/O)} : (N/b).B(R) + k.B(R)$$

2.4 Fourth I/O mechanism

2.4.1 Implementation

The fourth I/O mechanism needs to implement a read and write function performed by mapping and unmapping a B element portion of the file into internal memory through memory mapping. Every time the I/O mechanism needs to read and write outside of the of the mapped portion, the next B element of the file is mapped. In order to implement it the map method of the `java.nio.channels.FileChannel` class was used.

Memory mapped files are special, it allows Java program to access contents directly from memory, this is achieved by mapping whole file or portion of file into memory and operating system takes care of loading page requested and writing into file while application only deals with memory which results in very fast I/O operations. Memory used to load Memory mapped file is outside of Java heap Space. Java programming language supports memory mapped file with the `java.nio` package and has `MappedByteBuffer` to read and write from memory.

2.4.2 Expected behavior

Before running any test, a cost formula needs to be build for this implementation. The cost estimates the total number of I/O's that need to be done in function of N, B and k.

- k - the number of streams, where each stream is a file
- N - the size of the input file, measured in number of 32-bit integers.
- b - the size of the buffer in internal memory
- B(R) - the number of blocks that R occupies on disk

Cost(4th I/O) : $k \cdot B(R)$

2.5 Experimental observations

In order to perform experimentation over our streams classes, two benchmark were created. One benchmark allows to test all the read functions from the four different implementations, and the other one allows to test all the write functions. With the JMH benchmark library the program can apply test for different value of the same parameter and also do warm up before running real test. 2 thread are doing the test in order to get efficient results.

2.6 Configuration

The configuration of the computer that ran the test is the following :



Figure 1: Computer configuration




Figure 2: Free space on the computer

2.6.1 Write-Benchmark

So first the goal is to compare the implementations of each write functions. For the first two implementation of the I/O mechanism there is only two parameters to vary, N (size of the input file) and k the number of files. The two next implementation have one more parameter, B the size of the Buffer.

As the JMH benchmark library allow us the test will be done with various value for each parameters.

For the first run, k is taking its value in the following list $\{ "1", "2", "3", "4", "5", "10", "15", "20", "25", "30" \}$, and for each value of k a test is done with an input file of 100.000 in 32-bit integers and a buffer of size 1000 in 32-bits integers. In order to test the second implementation the size of the buffer is set to 8192 for the buffer implementation.

```

Result "com.ulb.psk.MyBenchmark.writeIntFile":
  7,132 ±(99.9%) 0,326 ms/op [Average]
  (min, avg, max) = (6,867, 7,132, 7,402), stdev = 0,216
  CI (99.9%): [6,805, 7,458] (assumes normal distribution)

# Run complete. Total time: 00:13:58

Benchmark (bufferSize) (dataVolume) (filesToWrite) Mode Cnt Score Error Units
MyBenchmark.bufferedWriteIntFile 1000 1000 1 avgt 10 0,034 ± 0,001 ms/op
MyBenchmark.bufferedWriteIntFile 1000 1000 2 avgt 10 0,040 ± 0,002 ms/op
MyBenchmark.bufferedWriteIntFile 1000 1000 3 avgt 10 0,042 ± 0,005 ms/op
MyBenchmark.bufferedWriteIntFile 1000 1000 4 avgt 10 0,045 ± 0,006 ms/op
MyBenchmark.bufferedWriteIntFile 1000 1000 5 avgt 10 0,045 ± 0,005 ms/op
MyBenchmark.bufferedWriteIntFile 1000 1000 10 avgt 10 0,045 ± 0,009 ms/op
MyBenchmark.bufferedWriteIntFile 1000 1000 15 avgt 10 0,046 ± 0,006 ms/op
MyBenchmark.bufferedWriteIntFile 1000 1000 20 avgt 10 0,047 ± 0,006 ms/op
MyBenchmark.bufferedWriteIntFile 1000 1000 25 avgt 10 0,037 ± 0,008 ms/op
MyBenchmark.bufferedWriteIntFile 1000 1000 30 avgt 10 0,032 ± 0,004 ms/op
MyBenchmark.bufferedWriteIntFile 8192 1000 1 avgt 10 0,039 ± 0,002 ms/op
MyBenchmark.bufferedWriteIntFile 8192 1000 2 avgt 10 0,035 ± 0,001 ms/op
MyBenchmark.bufferedWriteIntFile 8192 1000 3 avgt 10 0,033 ± 0,001 ms/op
MyBenchmark.bufferedWriteIntFile 8192 1000 4 avgt 10 0,031 ± 0,002 ms/op
MyBenchmark.bufferedWriteIntFile 8192 1000 5 avgt 10 0,031 ± 0,005 ms/op
MyBenchmark.bufferedWriteIntFile 8192 1000 10 avgt 10 0,032 ± 0,004 ms/op
MyBenchmark.bufferedWriteIntFile 8192 1000 15 avgt 10 0,031 ± 0,002 ms/op
MyBenchmark.bufferedWriteIntFile 8192 1000 20 avgt 10 0,031 ± 0,003 ms/op
MyBenchmark.bufferedWriteIntFile 8192 1000 25 avgt 10 0,031 ± 0,002 ms/op
MyBenchmark.bufferedWriteIntFile 8192 1000 30 avgt 10 0,032 ± 0,001 ms/op
MyBenchmark.mmWriteIntFile 1000 1000 1 avgt 10 0,018 ± 0,001 ms/op
MyBenchmark.mmWriteIntFile 1000 1000 2 avgt 10 0,025 ± 0,001 ms/op
MyBenchmark.mmWriteIntFile 1000 1000 3 avgt 10 0,031 ± 0,002 ms/op
MyBenchmark.mmWriteIntFile 1000 1000 4 avgt 10 0,037 ± 0,002 ms/op
MyBenchmark.mmWriteIntFile 1000 1000 5 avgt 10 0,042 ± 0,004 ms/op
MyBenchmark.mmWriteIntFile 1000 1000 10 avgt 10 0,067 ± 0,007 ms/op
MyBenchmark.mmWriteIntFile 1000 1000 15 avgt 10 0,094 ± 0,021 ms/op
MyBenchmark.mmWriteIntFile 1000 1000 20 avgt 10 0,121 ± 0,031 ms/op
MyBenchmark.mmWriteIntFile 1000 1000 25 avgt 10 0,149 ± 0,044 ms/op
MyBenchmark.mmWriteIntFile 1000 1000 30 avgt 10 0,153 ± 0,017 ms/op
MyBenchmark.writeIntFile N/A 1000 1 avgt 10 7,454 ± 0,526 ms/op
MyBenchmark.writeIntFile N/A 1000 2 avgt 10 7,577 ± 0,605 ms/op
MyBenchmark.writeIntFile N/A 1000 3 avgt 10 7,414 ± 0,405 ms/op
MyBenchmark.writeIntFile N/A 1000 4 avgt 10 7,230 ± 0,316 ms/op
MyBenchmark.writeIntFile N/A 1000 5 avgt 10 7,254 ± 0,327 ms/op
MyBenchmark.writeIntFile N/A 1000 10 avgt 10 7,080 ± 0,402 ms/op
MyBenchmark.writeIntFile N/A 1000 15 avgt 10 6,855 ± 0,473 ms/op
MyBenchmark.writeIntFile N/A 1000 20 avgt 10 6,398 ± 0,440 ms/op
MyBenchmark.writeIntFile N/A 1000 25 avgt 10 8,004 ± 0,876 ms/op
MyBenchmark.writeIntFile N/A 1000 30 avgt 10 7,132 ± 0,326 ms/op

```

Figure 3: Test N = 1.000

The first run reveals that the first write mechanism (WriteInt) is from far away the one that take the longest time. The write mechanism with the memory mapping seems to be slowest than the two implementation with the buffer. More over the second I/O mechanism (using a buffer of size 8192) is faster than the third I/O mechanism (using a buffer of size 1000). The best combination here is to use memory mapping with 1 file to write and a buffer of 1000 32-bits integers.

For the second run the value of the input file (N) is set to 100.000 32-bits integers. In order to test the second implementation the size of the buffer is set to 8192 for the buffer implementation.


```
Result "com.ulb.psk.MyBenchmark.writeIntFile":
576,078 ±(99.9%) 13,423 ms/op [Average]
(min, avg, max) = (556,532, 576,078, 585,086), stdev = 8,878
CI (99.9%): [562,655, 589,501] (assumes normal distribution)
```

```
# Run complete. Total time: 00:14:27
```

Benchmark	(bufferSize)	(dataVolume)	(filesToWrite)	Mode	Cnt	Score	Error	Units
MyBenchmark.bufferedWriteIntFile	8192	100000	1	avgt	10	3,415 ±	0,195	ms/op
MyBenchmark.bufferedWriteIntFile	8192	100000	2	avgt	10	3,375 ±	0,073	ms/op
MyBenchmark.bufferedWriteIntFile	8192	100000	3	avgt	10	3,162 ±	0,061	ms/op
MyBenchmark.bufferedWriteIntFile	8192	100000	4	avgt	10	3,273 ±	0,088	ms/op
MyBenchmark.bufferedWriteIntFile	8192	100000	5	avgt	10	2,790 ±	0,113	ms/op
MyBenchmark.bufferedWriteIntFile	8192	100000	10	avgt	10	3,021 ±	0,096	ms/op
MyBenchmark.bufferedWriteIntFile	8192	100000	15	avgt	10	2,413 ±	0,081	ms/op
MyBenchmark.bufferedWriteIntFile	8192	100000	20	avgt	10	2,972 ±	0,319	ms/op
MyBenchmark.bufferedWriteIntFile	8192	100000	25	avgt	10	2,684 ±	0,545	ms/op
MyBenchmark.bufferedWriteIntFile	8192	100000	30	avgt	10	2,441 ±	0,082	ms/op
MyBenchmark.bufferedWriteIntFile	10000	100000	1	avgt	10	3,331 ±	0,215	ms/op
MyBenchmark.bufferedWriteIntFile	10000	100000	2	avgt	10	3,252 ±	0,100	ms/op
MyBenchmark.bufferedWriteIntFile	10000	100000	3	avgt	10	3,431 ±	0,112	ms/op
MyBenchmark.bufferedWriteIntFile	10000	100000	4	avgt	10	3,444 ±	0,515	ms/op
MyBenchmark.bufferedWriteIntFile	10000	100000	5	avgt	10	2,793 ±	0,098	ms/op
MyBenchmark.bufferedWriteIntFile	10000	100000	10	avgt	10	2,504 ±	0,213	ms/op
MyBenchmark.bufferedWriteIntFile	10000	100000	15	avgt	10	2,311 ±	0,065	ms/op
MyBenchmark.bufferedWriteIntFile	10000	100000	20	avgt	10	2,311 ±	0,043	ms/op
MyBenchmark.bufferedWriteIntFile	10000	100000	25	avgt	10	2,292 ±	0,029	ms/op
MyBenchmark.bufferedWriteIntFile	10000	100000	30	avgt	10	2,363 ±	0,127	ms/op
MyBenchmark.mmWriteIntFile	10000	100000	1	avgt	10	1,463 ±	0,204	ms/op
MyBenchmark.mmWriteIntFile	10000	100000	2	avgt	10	1,700 ±	0,257	ms/op
MyBenchmark.mmWriteIntFile	10000	100000	3	avgt	10	1,575 ±	0,100	ms/op
MyBenchmark.mmWriteIntFile	10000	100000	4	avgt	10	1,579 ±	0,044	ms/op
MyBenchmark.mmWriteIntFile	10000	100000	5	avgt	10	1,514 ±	0,091	ms/op
MyBenchmark.mmWriteIntFile	10000	100000	10	avgt	10	1,768 ±	0,070	ms/op
MyBenchmark.mmWriteIntFile	10000	100000	15	avgt	10	1,931 ±	0,149	ms/op
MyBenchmark.mmWriteIntFile	10000	100000	20	avgt	10	2,415 ±	0,063	ms/op
MyBenchmark.mmWriteIntFile	10000	100000	25	avgt	10	2,521 ±	0,042	ms/op
MyBenchmark.mmWriteIntFile	10000	100000	30	avgt	10	2,637 ±	0,051	ms/op
MyBenchmark.writeIntFile	N/A	100000	1	avgt	10	666,993 ±	30,079	ms/op
MyBenchmark.writeIntFile	N/A	100000	2	avgt	10	658,243 ±	43,279	ms/op
MyBenchmark.writeIntFile	N/A	100000	3	avgt	10	534,820 ±	15,031	ms/op
MyBenchmark.writeIntFile	N/A	100000	4	avgt	10	546,347 ±	19,081	ms/op
MyBenchmark.writeIntFile	N/A	100000	5	avgt	10	549,830 ±	18,702	ms/op
MyBenchmark.writeIntFile	N/A	100000	10	avgt	10	531,809 ±	14,155	ms/op
MyBenchmark.writeIntFile	N/A	100000	15	avgt	10	551,015 ±	30,062	ms/op
MyBenchmark.writeIntFile	N/A	100000	20	avgt	10	547,186 ±	30,756	ms/op
MyBenchmark.writeIntFile	N/A	100000	25	avgt	10	569,422 ±	21,278	ms/op
MyBenchmark.writeIntFile	N/A	100000	30	avgt	10	576,078 ±	13,423	ms/op

Figure 4: Test N = 100.000

The results of the second run shows again that the writeInt function is the slowest. The best combination here is to use memory mapping with 1 file to write and a buffer of 1000 32-bits integers (same as the first run).

For the third run the value of the input file (N) is set to 1.000.000 32-bits integers. In order to test the second implementation the size of the buffer is set to 8192 for the buffer implementation.

```

Result "com.ulb.psk.MyBenchmark.writeIntFile":
5243,696 ±(99.9%) 36,783 ms/op [Average]
(min, avg, max) = (5221,530, 5243,696, 5288,271), stdev = 24,330
CI (99.9%): [5206,912, 5280,479] (assumes normal distribution)

```

```
# Run complete. Total time: 00:28:55
```

Benchmark	(bufferSize)	(dataVolume)	(filesToWrite)	Mode	Cnt	Score	Error	Units
MyBenchmark.bufferedWriteIntFile	1000	1000000	1	avgt	10	31,488 ±	2,583	ms/op
MyBenchmark.bufferedWriteIntFile	1000	1000000	2	avgt	10	31,625 ±	3,924	ms/op
MyBenchmark.bufferedWriteIntFile	1000	1000000	3	avgt	10	29,895 ±	2,061	ms/op
MyBenchmark.bufferedWriteIntFile	1000	1000000	4	avgt	10	30,093 ±	0,919	ms/op
MyBenchmark.bufferedWriteIntFile	1000	1000000	5	avgt	10	28,734 ±	1,137	ms/op
MyBenchmark.bufferedWriteIntFile	1000	1000000	10	avgt	10	28,576 ±	2,353	ms/op
MyBenchmark.bufferedWriteIntFile	1000	1000000	15	avgt	10	27,877 ±	1,394	ms/op
MyBenchmark.bufferedWriteIntFile	1000	1000000	20	avgt	10	37,313 ±	4,127	ms/op
MyBenchmark.bufferedWriteIntFile	1000	1000000	25	avgt	10	40,555 ±	5,169	ms/op
MyBenchmark.bufferedWriteIntFile	1000	1000000	30	avgt	10	38,937 ±	1,978	ms/op
MyBenchmark.bufferedWriteIntFile	8192	1000000	1	avgt	10	27,970 ±	1,387	ms/op
MyBenchmark.bufferedWriteIntFile	8192	1000000	2	avgt	10	28,949 ±	1,990	ms/op
MyBenchmark.bufferedWriteIntFile	8192	1000000	3	avgt	10	27,528 ±	0,733	ms/op
MyBenchmark.bufferedWriteIntFile	8192	1000000	4	avgt	10	30,770 ±	0,683	ms/op
MyBenchmark.bufferedWriteIntFile	8192	1000000	5	avgt	10	28,457 ±	4,273	ms/op
MyBenchmark.bufferedWriteIntFile	8192	1000000	10	avgt	10	28,249 ±	1,199	ms/op
MyBenchmark.bufferedWriteIntFile	8192	1000000	15	avgt	10	26,513 ±	1,031	ms/op
MyBenchmark.bufferedWriteIntFile	8192	1000000	20	avgt	10	26,397 ±	1,132	ms/op
MyBenchmark.bufferedWriteIntFile	8192	1000000	25	avgt	10	27,590 ±	2,264	ms/op
MyBenchmark.bufferedWriteIntFile	8192	1000000	30	avgt	10	28,284 ±	1,222	ms/op
MyBenchmark.mmWriteIntFile	1000	1000000	1	avgt	10	36,471 ±	6,052	ms/op
MyBenchmark.mmWriteIntFile	1000	1000000	2	avgt	10	58,679 ±	8,583	ms/op
MyBenchmark.mmWriteIntFile	1000	1000000	3	avgt	10	65,722 ±	4,653	ms/op
MyBenchmark.mmWriteIntFile	1000	1000000	4	avgt	10	71,769 ±	8,123	ms/op
MyBenchmark.mmWriteIntFile	1000	1000000	5	avgt	10	56,375 ±	6,772	ms/op
MyBenchmark.mmWriteIntFile	1000	1000000	10	avgt	10	57,165 ±	3,812	ms/op
MyBenchmark.mmWriteIntFile	1000	1000000	15	avgt	10	58,526 ±	4,676	ms/op
MyBenchmark.mmWriteIntFile	1000	1000000	20	avgt	10	61,822 ±	4,098	ms/op
MyBenchmark.mmWriteIntFile	1000	1000000	25	avgt	10	59,116 ±	5,044	ms/op
MyBenchmark.mmWriteIntFile	1000	1000000	30	avgt	10	61,558 ±	8,208	ms/op
MyBenchmark.writeIntFile	N/A	1000000	1	avgt	10	5676,961 ±	917,509	ms/op
MyBenchmark.writeIntFile	N/A	1000000	2	avgt	10	5195,480 ±	47,024	ms/op
MyBenchmark.writeIntFile	N/A	1000000	3	avgt	10	5228,343 ±	108,065	ms/op
MyBenchmark.writeIntFile	N/A	1000000	4	avgt	10	5198,916 ±	64,334	ms/op
MyBenchmark.writeIntFile	N/A	1000000	5	avgt	10	5407,696 ±	59,048	ms/op
MyBenchmark.writeIntFile	N/A	1000000	10	avgt	10	5214,788 ±	62,867	ms/op
MyBenchmark.writeIntFile	N/A	1000000	15	avgt	10	5211,016 ±	79,283	ms/op
MyBenchmark.writeIntFile	N/A	1000000	20	avgt	10	5252,365 ±	103,303	ms/op
MyBenchmark.writeIntFile	N/A	1000000	25	avgt	10	5310,124 ±	142,772	ms/op
MyBenchmark.writeIntFile	N/A	1000000	30	avgt	10	5243,696 ±	36,783	ms/op

Figure 5: Test N = 1.000.000

The third run shows again that the first I/O mechanism is very slow comparing to the other mechanisms. It tends to put away this implementation for the following of the project. This third brings something new, the memory mapping is not anymore the fastest. This time the best combination to use is to use buffered write with a buffer size of 8192 (second mechanism) and 20 files to write.

For the fourth run the value of the input file (N) is set to 5.000.000 32 bits integers. In order to test the second implementation the size of the buffer is set to 8192 for the buffer implementation.

```

Result "com.ulb.psk.MyBenchmark.writeIntFile":
26077,913 ±(99.9%) 266,902 ms/op [Average]
(min, avg, max) = (25833,601, 26077,913, 26374,008), stdev = 176,539
CI (99.9%): [25811,011, 26344,816] (assumes normal distribution)

```

```
# Run complete. Total time: 01:45:18
```

Benchmark	(bufferSize)	(dataVolume)	(filesToWrite)	Mode	Cnt	Score	Error	Units
MyBenchmark.bufferedWriteIntFile	1000	5000000	1	avgt	10	158,945 ±	10,611	ms/op
MyBenchmark.bufferedWriteIntFile	1000	5000000	2	avgt	10	186,286 ±	17,917	ms/op
MyBenchmark.bufferedWriteIntFile	1000	5000000	3	avgt	10	136,660 ±	6,781	ms/op
MyBenchmark.bufferedWriteIntFile	1000	5000000	4	avgt	10	137,560 ±	5,927	ms/op
MyBenchmark.bufferedWriteIntFile	1000	5000000	5	avgt	10	136,302 ±	3,211	ms/op
MyBenchmark.bufferedWriteIntFile	1000	5000000	10	avgt	10	137,602 ±	7,472	ms/op
MyBenchmark.bufferedWriteIntFile	1000	5000000	15	avgt	10	152,266 ±	11,195	ms/op
MyBenchmark.bufferedWriteIntFile	1000	5000000	20	avgt	10	170,588 ±	31,015	ms/op
MyBenchmark.bufferedWriteIntFile	1000	5000000	25	avgt	10	163,356 ±	6,142	ms/op
MyBenchmark.bufferedWriteIntFile	1000	5000000	30	avgt	10	157,701 ±	18,207	ms/op
MyBenchmark.bufferedWriteIntFile	8192	5000000	1	avgt	10	150,512 ±	11,115	ms/op
MyBenchmark.bufferedWriteIntFile	8192	5000000	2	avgt	10	176,720 ±	11,615	ms/op
MyBenchmark.bufferedWriteIntFile	8192	5000000	3	avgt	10	133,777 ±	13,793	ms/op
MyBenchmark.bufferedWriteIntFile	8192	5000000	4	avgt	10	135,368 ±	14,459	ms/op
MyBenchmark.bufferedWriteIntFile	8192	5000000	5	avgt	10	132,724 ±	22,567	ms/op
MyBenchmark.bufferedWriteIntFile	8192	5000000	10	avgt	10	141,408 ±	4,997	ms/op
MyBenchmark.bufferedWriteIntFile	8192	5000000	15	avgt	10	131,879 ±	21,411	ms/op
MyBenchmark.bufferedWriteIntFile	8192	5000000	20	avgt	10	133,963 ±	5,649	ms/op
MyBenchmark.bufferedWriteIntFile	8192	5000000	25	avgt	10	133,665 ±	3,733	ms/op
MyBenchmark.bufferedWriteIntFile	8192	5000000	30	avgt	10	134,609 ±	5,500	ms/op
MyBenchmark.mmWriteIntFile	1000	5000000	1	avgt	10	218,697 ±	30,948	ms/op
MyBenchmark.mmWriteIntFile	1000	5000000	2	avgt	10	309,910 ±	29,738	ms/op
MyBenchmark.mmWriteIntFile	1000	5000000	3	avgt	10	259,758 ±	9,885	ms/op
MyBenchmark.mmWriteIntFile	1000	5000000	4	avgt	10	254,948 ±	8,377	ms/op
MyBenchmark.mmWriteIntFile	1000	5000000	5	avgt	10	256,975 ±	11,229	ms/op
MyBenchmark.mmWriteIntFile	1000	5000000	10	avgt	10	268,300 ±	10,107	ms/op
MyBenchmark.mmWriteIntFile	1000	5000000	15	avgt	10	287,143 ±	73,404	ms/op
MyBenchmark.mmWriteIntFile	1000	5000000	20	avgt	10	284,672 ±	59,699	ms/op
MyBenchmark.mmWriteIntFile	1000	5000000	25	avgt	10	280,410 ±	20,309	ms/op
MyBenchmark.mmWriteIntFile	1000	5000000	30	avgt	10	283,780 ±	17,110	ms/op
MyBenchmark.writeIntFile	N/A	5000000	1	avgt	10	29136,875 ±	3002,747	ms/op
MyBenchmark.writeIntFile	N/A	5000000	2	avgt	10	29273,168 ±	642,624	ms/op
MyBenchmark.writeIntFile	N/A	5000000	3	avgt	10	28798,192 ±	856,384	ms/op
MyBenchmark.writeIntFile	N/A	5000000	4	avgt	10	29360,537 ±	953,454	ms/op
MyBenchmark.writeIntFile	N/A	5000000	5	avgt	10	26579,117 ±	978,690	ms/op
MyBenchmark.writeIntFile	N/A	5000000	10	avgt	10	25999,583 ±	293,830	ms/op
MyBenchmark.writeIntFile	N/A	5000000	15	avgt	10	26013,073 ±	409,329	ms/op
MyBenchmark.writeIntFile	N/A	5000000	20	avgt	10	25926,156 ±	161,744	ms/op
MyBenchmark.writeIntFile	N/A	5000000	25	avgt	10	26040,508 ±	254,991	ms/op
MyBenchmark.writeIntFile	N/A	5000000	30	avgt	10	26077,913 ±	266,902	ms/op

Figure 6: Test N = 5.000.000

The fourth run as the previous run is showing that the write int file is obsolete comparing to the other implementation. According to this run the best combination is to use a buffered write int with a buffer size of 8192 (second mechanism) and 15 files to write.

For the fifth run the value of the input file (N) is set to 7.000.000 32 bits integers. In order to test the second implementation the size of the buffer is set to 8192 for the buffer implementation.

```
Result "com.ulb.psk.MyBenchmark.writeIntFile":
37479,286 ±(99.9%) 272,799 ms/op [Average]
(min, avg, max) = (37190,127, 37479,286, 37739,204), stdev = 180,439
CI (99.9%): [37206,487, 37752,084] (assumes normal distribution)
```

```
# Run complete. Total time: 02:23:25
```

Benchmark	(bufferSize)	(dataVolume)	(filesToWrite)	Mode	Cnt	Score	Error	Units
MyBenchmark.bufferedWriteIntFile	1000	7000000	1	avgt	10	246,602 ±	25,101	ms/op
MyBenchmark.bufferedWriteIntFile	1000	7000000	2	avgt	10	343,961 ±	54,897	ms/op
MyBenchmark.bufferedWriteIntFile	1000	7000000	3	avgt	10	280,213 ±	35,998	ms/op
MyBenchmark.bufferedWriteIntFile	1000	7000000	4	avgt	10	226,426 ±	43,366	ms/op
MyBenchmark.bufferedWriteIntFile	1000	7000000	5	avgt	10	208,724 ±	12,784	ms/op
MyBenchmark.bufferedWriteIntFile	1000	7000000	10	avgt	10	238,669 ±	79,697	ms/op
MyBenchmark.bufferedWriteIntFile	1000	7000000	15	avgt	10	241,603 ±	32,532	ms/op
MyBenchmark.bufferedWriteIntFile	1000	7000000	20	avgt	10	243,502 ±	16,605	ms/op
MyBenchmark.bufferedWriteIntFile	1000	7000000	25	avgt	10	243,055 ±	58,140	ms/op
MyBenchmark.bufferedWriteIntFile	1000	7000000	30	avgt	10	222,433 ±	8,939	ms/op
MyBenchmark.bufferedWriteIntFile	8192	7000000	1	avgt	10	213,156 ±	34,732	ms/op
MyBenchmark.bufferedWriteIntFile	8192	7000000	2	avgt	10	215,379 ±	11,878	ms/op
MyBenchmark.bufferedWriteIntFile	8192	7000000	3	avgt	10	225,056 ±	8,466	ms/op
MyBenchmark.bufferedWriteIntFile	8192	7000000	4	avgt	10	179,298 ±	15,211	ms/op
MyBenchmark.bufferedWriteIntFile	8192	7000000	5	avgt	10	174,631 ±	16,059	ms/op
MyBenchmark.bufferedWriteIntFile	8192	7000000	10	avgt	10	190,051 ±	17,216	ms/op
MyBenchmark.bufferedWriteIntFile	8192	7000000	15	avgt	10	195,374 ±	94,714	ms/op
MyBenchmark.bufferedWriteIntFile	8192	7000000	20	avgt	10	174,947 ±	12,949	ms/op
MyBenchmark.bufferedWriteIntFile	8192	7000000	25	avgt	10	174,130 ±	5,611	ms/op
MyBenchmark.bufferedWriteIntFile	8192	7000000	30	avgt	10	175,708 ±	5,011	ms/op
MyBenchmark.mmWriteIntFile	1000	7000000	1	avgt	10	322,388 ±	71,113	ms/op
MyBenchmark.mmWriteIntFile	1000	7000000	2	avgt	10	477,561 ±	52,187	ms/op
MyBenchmark.mmWriteIntFile	1000	7000000	3	avgt	10	403,982 ±	29,761	ms/op
MyBenchmark.mmWriteIntFile	1000	7000000	4	avgt	10	360,681 ±	15,064	ms/op
MyBenchmark.mmWriteIntFile	1000	7000000	5	avgt	10	367,399 ±	18,467	ms/op
MyBenchmark.mmWriteIntFile	1000	7000000	10	avgt	10	383,425 ±	67,196	ms/op
MyBenchmark.mmWriteIntFile	1000	7000000	15	avgt	10	375,352 ±	14,020	ms/op
MyBenchmark.mmWriteIntFile	1000	7000000	20	avgt	10	387,260 ±	57,010	ms/op
MyBenchmark.mmWriteIntFile	1000	7000000	25	avgt	10	438,789 ±	316,200	ms/op
MyBenchmark.mmWriteIntFile	1000	7000000	30	avgt	10	400,781 ±	81,421	ms/op
MyBenchmark.writeIntFile	N/A	7000000	1	avgt	10	38779,091 ±	407,859	ms/op
MyBenchmark.writeIntFile	N/A	7000000	2	avgt	10	39088,631 ±	825,534	ms/op
MyBenchmark.writeIntFile	N/A	7000000	3	avgt	10	38418,280 ±	940,858	ms/op
MyBenchmark.writeIntFile	N/A	7000000	4	avgt	10	38389,754 ±	335,040	ms/op
MyBenchmark.writeIntFile	N/A	7000000	5	avgt	10	38349,884 ±	237,246	ms/op
MyBenchmark.writeIntFile	N/A	7000000	10	avgt	10	38312,227 ±	322,453	ms/op
MyBenchmark.writeIntFile	N/A	7000000	15	avgt	10	38637,758 ±	368,924	ms/op
MyBenchmark.writeIntFile	N/A	7000000	20	avgt	10	38333,828 ±	587,904	ms/op
MyBenchmark.writeIntFile	N/A	7000000	25	avgt	10	37837,173 ±	275,679	ms/op
MyBenchmark.writeIntFile	N/A	7000000	30	avgt	10	37479,286 ±	272,799	ms/op

Figure 7: Test N = 7.000.000

The best combination for the fifth run is to use a buffered write int with a buffer size of 8192 (second mechanism) and 25 files to write.

The value of N is becoming too big to keep on doing test on the machine, because of the free memory. So it will stop the test on varying the parameter N.

Now that different value of N and k have been tried there is just the size of the buffer to play with. In order to test the impact of the buffer size, B is set to multiple value; 8192, 10.000, 50.000, 100.000. More over in order to test this with a big size in input, N is set to 10.000.000 32-bits integers. This lead to the following results :

Result "com.ulb.psk.MyBenchmark.writeIntFile":
 53167,760 ±(99.9%) 1679,723 ms/op [Average]
 (min, avg, max) = (51077,540, 53167,760, 54485,088), stdev = 1111,033
 CI (99.9%): [51488,037, 54847,483] (assumes normal distribution)

Run complete. Total time: 03:32:55

Benchmark	(bufferSize)	(dataVolume)	(filesToWrite)	Mode	Cnt	Score	Error	Units
MyBenchmark.bufferedWriteIntFile	8192	10000000	1	avgt	10	295,083 ±	32,498	ms/op
MyBenchmark.bufferedWriteIntFile	8192	10000000	2	avgt	10	328,281 ±	20,370	ms/op
MyBenchmark.bufferedWriteIntFile	8192	10000000	3	avgt	10	299,121 ±	17,597	ms/op
MyBenchmark.bufferedWriteIntFile	8192	10000000	4	avgt	10	329,472 ±	21,563	ms/op
MyBenchmark.bufferedWriteIntFile	8192	10000000	5	avgt	10	239,296 ±	8,482	ms/op
MyBenchmark.bufferedWriteIntFile	8192	10000000	10	avgt	10	250,893 ±	8,840	ms/op
MyBenchmark.bufferedWriteIntFile	8192	10000000	15	avgt	10	234,207 ±	10,635	ms/op
MyBenchmark.bufferedWriteIntFile	8192	10000000	20	avgt	10	267,195 ±	11,562	ms/op
MyBenchmark.bufferedWriteIntFile	8192	10000000	25	avgt	10	288,307 ±	35,548	ms/op
MyBenchmark.bufferedWriteIntFile	8192	10000000	30	avgt	10	282,037 ±	49,212	ms/op
MyBenchmark.bufferedWriteIntFile	10000	10000000	1	avgt	10	404,009 ±	109,481	ms/op
MyBenchmark.bufferedWriteIntFile	10000	10000000	2	avgt	10	409,915 ±	82,194	ms/op
MyBenchmark.bufferedWriteIntFile	10000	10000000	3	avgt	10	363,684 ±	46,481	ms/op
MyBenchmark.bufferedWriteIntFile	10000	10000000	4	avgt	10	336,119 ±	45,109	ms/op
MyBenchmark.bufferedWriteIntFile	10000	10000000	5	avgt	10	226,898 ±	5,795	ms/op
MyBenchmark.bufferedWriteIntFile	10000	10000000	10	avgt	10	247,600 ±	8,024	ms/op
MyBenchmark.bufferedWriteIntFile	10000	10000000	15	avgt	10	228,914 ±	10,438	ms/op
MyBenchmark.bufferedWriteIntFile	10000	10000000	20	avgt	10	249,169 ±	20,951	ms/op
MyBenchmark.bufferedWriteIntFile	10000	10000000	25	avgt	10	251,125 ±	15,037	ms/op
MyBenchmark.bufferedWriteIntFile	10000	10000000	30	avgt	10	249,565 ±	13,146	ms/op
MyBenchmark.bufferedWriteIntFile	50000	10000000	1	avgt	10	460,588 ±	72,573	ms/op
MyBenchmark.bufferedWriteIntFile	50000	10000000	2	avgt	10	448,560 ±	53,377	ms/op
MyBenchmark.bufferedWriteIntFile	50000	10000000	3	avgt	10	396,207 ±	66,057	ms/op
MyBenchmark.bufferedWriteIntFile	50000	10000000	4	avgt	10	394,051 ±	37,873	ms/op
MyBenchmark.bufferedWriteIntFile	50000	10000000	5	avgt	10	233,242 ±	6,785	ms/op
MyBenchmark.bufferedWriteIntFile	50000	10000000	10	avgt	10	232,001 ±	8,559	ms/op
MyBenchmark.bufferedWriteIntFile	50000	10000000	15	avgt	10	251,956 ±	15,347	ms/op
MyBenchmark.bufferedWriteIntFile	50000	10000000	20	avgt	10	263,364 ±	22,866	ms/op
MyBenchmark.bufferedWriteIntFile	50000	10000000	25	avgt	10	265,834 ±	21,504	ms/op
MyBenchmark.bufferedWriteIntFile	50000	10000000	30	avgt	10	264,477 ±	23,185	ms/op
MyBenchmark.bufferedWriteIntFile	100000	10000000	1	avgt	10	673,708 ±	37,122	ms/op
MyBenchmark.bufferedWriteIntFile	100000	10000000	2	avgt	10	576,233 ±	60,068	ms/op
MyBenchmark.bufferedWriteIntFile	100000	10000000	3	avgt	10	511,549 ±	30,451	ms/op
MyBenchmark.bufferedWriteIntFile	100000	10000000	4	avgt	10	422,033 ±	44,640	ms/op
MyBenchmark.bufferedWriteIntFile	100000	10000000	5	avgt	10	241,867 ±	15,956	ms/op
MyBenchmark.bufferedWriteIntFile	100000	10000000	10	avgt	10	257,084 ±	76,225	ms/op
MyBenchmark.bufferedWriteIntFile	100000	10000000	15	avgt	10	251,288 ±	13,291	ms/op
MyBenchmark.bufferedWriteIntFile	100000	10000000	20	avgt	10	265,178 ±	25,387	ms/op
MyBenchmark.bufferedWriteIntFile	100000	10000000	25	avgt	10	258,564 ±	38,020	ms/op
MyBenchmark.bufferedWriteIntFile	100000	10000000	30	avgt	10	253,875 ±	20,630	ms/op
MyBenchmark.mmWriteIntFile	10000	10000000	1	avgt	10	141,667 ±	18,230	ms/op
MyBenchmark.mmWriteIntFile	10000	10000000	2	avgt	10	148,099 ±	16,673	ms/op
MyBenchmark.mmWriteIntFile	10000	10000000	3	avgt	10	153,998 ±	15,331	ms/op
MyBenchmark.mmWriteIntFile	10000	10000000	4	avgt	10	171,418 ±	33,659	ms/op

MyBenchmark.mmWriteIntFile	10000	10000000	5	avgt	10	148,646 ±	4,586	ms/op
MyBenchmark.mmWriteIntFile	10000	10000000	10	avgt	10	198,393 ±	24,786	ms/op
MyBenchmark.mmWriteIntFile	10000	10000000	15	avgt	10	197,663 ±	15,479	ms/op
MyBenchmark.mmWriteIntFile	10000	10000000	20	avgt	10	193,992 ±	9,119	ms/op
MyBenchmark.mmWriteIntFile	10000	10000000	25	avgt	10	201,246 ±	15,086	ms/op
MyBenchmark.mmWriteIntFile	10000	10000000	30	avgt	10	201,055 ±	6,947	ms/op
MyBenchmark.mmWriteIntFile	50000	10000000	1	avgt	10	126,795 ±	33,164	ms/op
MyBenchmark.mmWriteIntFile	50000	10000000	2	avgt	10	124,025 ±	14,798	ms/op
MyBenchmark.mmWriteIntFile	50000	10000000	3	avgt	10	122,004 ±	4,075	ms/op
MyBenchmark.mmWriteIntFile	50000	10000000	4	avgt	10	132,651 ±	17,042	ms/op
MyBenchmark.mmWriteIntFile	50000	10000000	5	avgt	10	128,199 ±	4,210	ms/op
MyBenchmark.mmWriteIntFile	50000	10000000	10	avgt	10	189,370 ±	5,171	ms/op
MyBenchmark.mmWriteIntFile	50000	10000000	15	avgt	10	213,757 ±	7,496	ms/op
MyBenchmark.mmWriteIntFile	50000	10000000	20	avgt	10	211,339 ±	6,367	ms/op
MyBenchmark.mmWriteIntFile	50000	10000000	25	avgt	10	212,452 ±	7,987	ms/op
MyBenchmark.mmWriteIntFile	50000	10000000	30	avgt	10	212,879 ±	12,983	ms/op
MyBenchmark.mmWriteIntFile	100000	10000000	1	avgt	10	114,845 ±	12,983	ms/op
MyBenchmark.mmWriteIntFile	100000	10000000	2	avgt	10	122,372 ±	9,270	ms/op
MyBenchmark.mmWriteIntFile	100000	10000000	3	avgt	10	121,497 ±	3,754	ms/op
MyBenchmark.mmWriteIntFile	100000	10000000	4	avgt	10	120,861 ±	5,530	ms/op
MyBenchmark.mmWriteIntFile	100000	10000000	5	avgt	10	124,593 ±	4,297	ms/op
MyBenchmark.mmWriteIntFile	100000	10000000	10	avgt	10	197,331 ±	7,427	ms/op
MyBenchmark.mmWriteIntFile	100000	10000000	15	avgt	10	223,016 ±	4,756	ms/op
MyBenchmark.mmWriteIntFile	100000	10000000	20	avgt	10	224,349 ±	5,632	ms/op
MyBenchmark.mmWriteIntFile	100000	10000000	25	avgt	10	230,758 ±	8,388	ms/op
MyBenchmark.mmWriteIntFile	100000	10000000	30	avgt	10	234,787 ±	17,118	ms/op
MyBenchmark.writeIntFile	N/A	10000000	1	avgt	10	54629,943 ±	443,312	ms/op
MyBenchmark.writeIntFile	N/A	10000000	2	avgt	10	54446,142 ±	267,360	ms/op
MyBenchmark.writeIntFile	N/A	10000000	3	avgt	10	54574,867 ±	530,059	ms/op
MyBenchmark.writeIntFile	N/A	10000000	4	avgt	10	54809,896 ±	842,873	ms/op
MyBenchmark.writeIntFile	N/A	10000000	5	avgt	10	55622,501 ±	1685,965	ms/op
MyBenchmark.writeIntFile	N/A	10000000	10	avgt	10	58722,711 ±	5686,321	ms/op
MyBenchmark.writeIntFile	N/A	10000000	15	avgt	10	55434,278 ±	645,896	ms/op
MyBenchmark.writeIntFile	N/A	10000000	20	avgt	10	53263,151 ±	2940,028	ms/op
MyBenchmark.writeIntFile	N/A	10000000	25	avgt	10	51678,350 ±	1651,657	ms/op
MyBenchmark.writeIntFile	N/A	10000000	30	avgt	10	53167,760 ±	1679,723	ms/op

Figure 8: 1st Test on buffer size

According to the previous run the buffered write int should be the best function to use here. But according to the result the best combination is to use the memory mapping with a buffer of 100.000 32-bits integers and 1 file to write.

In order to confirm the previous result, the data volume is set to 20.000.000 32-bits integers. More over the size of the buffer is also changed, now it is going to 1.000.000 and 2.000.000 32-bits integers.

Run complete. Total time: 01:00:20

Benchmark	(bufferSize)	(dataVolume)	(filesToWrite)	Mode	Cnt	Score	Error	Units
MyBenchmark.bufferedWriteIntFile	8192	20000000	1	avgt	10	511,067 ±	54,905	ms/op
MyBenchmark.bufferedWriteIntFile	8192	20000000	2	avgt	10	570,568 ±	51,070	ms/op
MyBenchmark.bufferedWriteIntFile	8192	20000000	3	avgt	10	628,796 ±	70,888	ms/op
MyBenchmark.bufferedWriteIntFile	8192	20000000	4	avgt	10	617,469 ±	53,087	ms/op
MyBenchmark.bufferedWriteIntFile	8192	20000000	5	avgt	10	646,213 ±	141,849	ms/op
MyBenchmark.bufferedWriteIntFile	8192	20000000	10	avgt	10	528,088 ±	24,170	ms/op
MyBenchmark.bufferedWriteIntFile	8192	20000000	15	avgt	10	476,849 ±	36,255	ms/op
MyBenchmark.bufferedWriteIntFile	8192	20000000	20	avgt	10	493,737 ±	19,791	ms/op
MyBenchmark.bufferedWriteIntFile	8192	20000000	25	avgt	10	530,713 ±	56,517	ms/op
MyBenchmark.bufferedWriteIntFile	8192	20000000	30	avgt	10	538,909 ±	89,318	ms/op
MyBenchmark.bufferedWriteIntFile	10000	20000000	1	avgt	10	1276,630 ±	609,513	ms/op
MyBenchmark.bufferedWriteIntFile	10000	20000000	2	avgt	10	1091,816 ±	411,150	ms/op
MyBenchmark.bufferedWriteIntFile	10000	20000000	3	avgt	10	1194,743 ±	462,344	ms/op
MyBenchmark.bufferedWriteIntFile	10000	20000000	4	avgt	10	1026,204 ±	409,809	ms/op
MyBenchmark.bufferedWriteIntFile	10000	20000000	5	avgt	10	1192,772 ±	263,295	ms/op
MyBenchmark.bufferedWriteIntFile	10000	20000000	10	avgt	10	548,092 ±	77,790	ms/op
MyBenchmark.bufferedWriteIntFile	10000	20000000	15	avgt	10	526,413 ±	157,298	ms/op
MyBenchmark.bufferedWriteIntFile	10000	20000000	20	avgt	10	541,929 ±	64,138	ms/op
MyBenchmark.bufferedWriteIntFile	10000	20000000	25	avgt	10	551,159 ±	44,987	ms/op
MyBenchmark.bufferedWriteIntFile	10000	20000000	30	avgt	10	607,091 ±	97,527	ms/op
MyBenchmark.bufferedWriteIntFile	100000	20000000	1	avgt	10	1471,175 ±	217,979	ms/op
MyBenchmark.bufferedWriteIntFile	100000	20000000	2	avgt	10	1309,037 ±	144,822	ms/op
MyBenchmark.bufferedWriteIntFile	100000	20000000	3	avgt	10	1274,185 ±	217,580	ms/op
MyBenchmark.bufferedWriteIntFile	100000	20000000	4	avgt	10	1152,132 ±	119,794	ms/op
MyBenchmark.bufferedWriteIntFile	100000	20000000	5	avgt	10	992,087 ±	169,667	ms/op
MyBenchmark.bufferedWriteIntFile	100000	20000000	10	avgt	10	514,825 ±	53,641	ms/op
MyBenchmark.bufferedWriteIntFile	100000	20000000	15	avgt	10	504,100 ±	38,010	ms/op
MyBenchmark.bufferedWriteIntFile	100000	20000000	20	avgt	10	529,546 ±	44,947	ms/op
MyBenchmark.bufferedWriteIntFile	100000	20000000	25	avgt	10	615,838 ±	299,228	ms/op
MyBenchmark.bufferedWriteIntFile	100000	20000000	30	avgt	10	559,687 ±	75,753	ms/op
MyBenchmark.bufferedWriteIntFile	1000000	20000000	1	avgt	10	1054,096 ±	87,645	ms/op
MyBenchmark.bufferedWriteIntFile	1000000	20000000	2	avgt	10	931,863 ±	100,612	ms/op
MyBenchmark.bufferedWriteIntFile	1000000	20000000	3	avgt	10	962,935 ±	114,415	ms/op
MyBenchmark.bufferedWriteIntFile	1000000	20000000	4	avgt	10	888,156 ±	57,201	ms/op
MyBenchmark.bufferedWriteIntFile	1000000	20000000	5	avgt	10	821,134 ±	42,465	ms/op
MyBenchmark.bufferedWriteIntFile	1000000	20000000	10	avgt	10	673,650 ±	50,082	ms/op
MyBenchmark.bufferedWriteIntFile	1000000	20000000	15	avgt	10	799,404 ±	65,330	ms/op
MyBenchmark.bufferedWriteIntFile	1000000	20000000	20	avgt	10	478,639 ±	52,549	ms/op
MyBenchmark.bufferedWriteIntFile	1000000	20000000	25	avgt	10	511,141 ±	121,880	ms/op
MyBenchmark.bufferedWriteIntFile	1000000	20000000	30	avgt	10	502,656 ±	53,691	ms/op
MyBenchmark.bufferedWriteIntFile	10000000	20000000	1	avgt	10	608,669 ±	16,118	ms/op
MyBenchmark.bufferedWriteIntFile	10000000	20000000	2	avgt	10	472,286 ±	19,499	ms/op
MyBenchmark.bufferedWriteIntFile	10000000	20000000	3	avgt	10	488,913 ±	36,866	ms/op
MyBenchmark.bufferedWriteIntFile	10000000	20000000	4	avgt	10	496,675 ±	15,782	ms/op
MyBenchmark.bufferedWriteIntFile	10000000	20000000	5	avgt	10	497,039 ±	36,294	ms/op
MyBenchmark.bufferedWriteIntFile	10000000	20000000	10	avgt	10	536,857 ±	88,897	ms/op
MyBenchmark.bufferedWriteIntFile	10000000	20000000	15	avgt	10	598,901 ±	326,812	ms/op
MyBenchmark.bufferedWriteIntFile	10000000	20000000	20	avgt	10	557,664 ±	117,861	ms/op
MyBenchmark.bufferedWriteIntFile	10000000	20000000	25	avgt	10	514,869 ±	142,716	ms/op
MyBenchmark.bufferedWriteIntFile	10000000	20000000	30	avgt	10	585,119 ±	173,860	ms/op
MyBenchmark.bufferedWriteIntFile	20000000	20000000	1	avgt	10	447,412 ±	32,828	ms/op
MyBenchmark.bufferedWriteIntFile	20000000	20000000	2	avgt	10	471,420 ±	51,054	ms/op
MyBenchmark.bufferedWriteIntFile	20000000	20000000	3	avgt	10	506,761 ±	83,750	ms/op
MyBenchmark.bufferedWriteIntFile	20000000	20000000	4	avgt	10	498,852 ±	86,463	ms/op
MyBenchmark.bufferedWriteIntFile	20000000	20000000	5	avgt	10	552,249 ±	88,789	ms/op

MyBenchmark.bufferedWriteIntFile	20000000	20000000	5	avgt	10	552,249	±	88,789	ms/op
MyBenchmark.bufferedWriteIntFile	20000000	20000000	10	avgt	10	545,704	±	144,435	ms/op
MyBenchmark.bufferedWriteIntFile	20000000	20000000	15	avgt	10	531,927	±	139,170	ms/op
MyBenchmark.mmWriteIntFile	10000	20000000	1	avgt	10	270,563	±	20,929	ms/op
MyBenchmark.mmWriteIntFile	10000	20000000	2	avgt	10	287,112	±	17,350	ms/op
MyBenchmark.mmWriteIntFile	10000	20000000	3	avgt	10	285,064	±	18,533	ms/op
MyBenchmark.mmWriteIntFile	10000	20000000	4	avgt	10	288,403	±	13,121	ms/op
MyBenchmark.mmWriteIntFile	10000	20000000	5	avgt	10	285,361	±	5,034	ms/op
MyBenchmark.mmWriteIntFile	10000	20000000	10	avgt	10	367,867	±	9,784	ms/op
MyBenchmark.mmWriteIntFile	10000	20000000	15	avgt	10	381,859	±	10,909	ms/op
MyBenchmark.mmWriteIntFile	10000	20000000	20	avgt	10	384,682	±	10,339	ms/op
MyBenchmark.mmWriteIntFile	10000	20000000	25	avgt	10	402,493	±	25,728	ms/op
MyBenchmark.mmWriteIntFile	10000	20000000	30	avgt	10	399,268	±	11,353	ms/op
MyBenchmark.mmWriteIntFile	100000	20000000	1	avgt	10	253,824	±	63,919	ms/op
MyBenchmark.mmWriteIntFile	100000	20000000	2	avgt	10	271,852	±	28,070	ms/op
MyBenchmark.mmWriteIntFile	100000	20000000	3	avgt	10	267,149	±	17,437	ms/op
MyBenchmark.mmWriteIntFile	100000	20000000	4	avgt	10	309,546	±	181,216	ms/op
MyBenchmark.mmWriteIntFile	100000	20000000	5	avgt	10	273,118	±	14,364	ms/op
MyBenchmark.mmWriteIntFile	100000	20000000	10	avgt	10	396,759	±	14,601	ms/op
MyBenchmark.mmWriteIntFile	100000	20000000	15	avgt	10	449,732	±	5,973	ms/op
MyBenchmark.mmWriteIntFile	100000	20000000	20	avgt	10	462,545	±	32,370	ms/op
MyBenchmark.mmWriteIntFile	100000	20000000	25	avgt	10	464,518	±	21,901	ms/op
MyBenchmark.mmWriteIntFile	100000	20000000	30	avgt	10	462,578	±	15,651	ms/op
MyBenchmark.mmWriteIntFile	1000000	20000000	1	avgt	10	589,498	±	60,801	ms/op
MyBenchmark.mmWriteIntFile	1000000	20000000	2	avgt	10	628,196	±	72,914	ms/op
MyBenchmark.mmWriteIntFile	1000000	20000000	3	avgt	10	574,631	±	79,179	ms/op
MyBenchmark.mmWriteIntFile	1000000	20000000	4	avgt	10	543,273	±	63,791	ms/op
MyBenchmark.mmWriteIntFile	1000000	20000000	5	avgt	10	510,464	±	52,149	ms/op
MyBenchmark.mmWriteIntFile	1000000	20000000	10	avgt	10	579,650	±	43,063	ms/op
MyBenchmark.mmWriteIntFile	1000000	20000000	15	avgt	10	729,632	±	28,133	ms/op
MyBenchmark.mmWriteIntFile	1000000	20000000	20	avgt	10	446,097	±	20,726	ms/op
MyBenchmark.mmWriteIntFile	1000000	20000000	25	avgt	10	455,381	±	18,740	ms/op
MyBenchmark.mmWriteIntFile	1000000	20000000	30	avgt	10	472,575	±	24,056	ms/op
MyBenchmark.mmWriteIntFile	10000000	20000000	1	avgt	10	236,929	±	12,579	ms/op
MyBenchmark.mmWriteIntFile	10000000	20000000	2	avgt	10	219,880	±	16,804	ms/op
MyBenchmark.mmWriteIntFile	10000000	20000000	3	avgt	10	213,703	±	18,547	ms/op
MyBenchmark.mmWriteIntFile	10000000	20000000	4	avgt	10	208,655	±	10,678	ms/op
MyBenchmark.mmWriteIntFile	10000000	20000000	5	avgt	10	209,069	±	13,405	ms/op
MyBenchmark.mmWriteIntFile	10000000	20000000	10	avgt	10	365,499	±	4,485	ms/op
MyBenchmark.mmWriteIntFile	10000000	20000000	15	avgt	8	459,802	±	75,592	ms/op
MyBenchmark.mmWriteIntFile	10000000	20000000	20	avgt	5	468,664	±	37,525	ms/op
MyBenchmark.mmWriteIntFile	10000000	20000000	25	avgt	4	501,647	±	206,111	ms/op
MyBenchmark.mmWriteIntFile	10000000	20000000	30	avgt	6	529,865	±	70,105	ms/op
MyBenchmark.mmWriteIntFile	20000000	20000000	1	avgt	10	262,690	±	81,253	ms/op
MyBenchmark.mmWriteIntFile	20000000	20000000	2	avgt	10	264,602	±	28,818	ms/op
MyBenchmark.mmWriteIntFile	20000000	20000000	3	avgt	10	265,337	±	26,037	ms/op
MyBenchmark.mmWriteIntFile	20000000	20000000	4	avgt	10	242,837	±	30,159	ms/op
MyBenchmark.mmWriteIntFile	20000000	20000000	5	avgt	6	286,358	±	111,618	ms/op
MyBenchmark.mmWriteIntFile	20000000	20000000	10	avgt	5	420,907	±	69,193	ms/op
MyBenchmark.mmWriteIntFile	20000000	20000000	15	avgt	3	471,738	±	372,433	ms/op

Figure 9: 2nd Test on buffer size

All of those previous benchmarks leads to the fact that the memory mapping mechanisms is the fastest and it doesn't required a lot of file in input in order to get a proper score.

2.6.2 Read-Benchmark

Now that all the implementation of the write function are done. The focus is done on the different implementations of the read function. After running a few benchmarks, it directly appears that once more the fourth mechanisms is the fastest and that the second and third one have almost the same time of execution.

For the first run, k is taking its value in the following list {"1","2","3","4","5","10","15","20","25","30"}, and for each value of k a test is done with an input file of 100.000 in 32-bit integers and a buffer of size 1.000 in 32-bits Integers.

```
Result "com.ulb.psk.MyBenchmark.readIntFile":
3,515 ±(99.9%) 0,149 ms/op [Average]
(min, avg, max) = (3,406, 3,515, 3,712), stdev = 0,099
CI (99.9%): [3,366, 3,664] (assumes normal distribution)

# Run complete. Total time: 00:07:13
```

Benchmark	(bufferSize)	(dataVolume)	(filesToRead)	Mode	Cnt	Score	Error	Units
MyBenchmark.mmReadIntFile	1000	1000	1	avgt	10	0,019 ± 0,003		ms/op
MyBenchmark.mmReadIntFile	1000	1000	2	avgt	10	0,022 ± 0,004		ms/op
MyBenchmark.mmReadIntFile	1000	1000	3	avgt	10	0,024 ± 0,004		ms/op
MyBenchmark.mmReadIntFile	1000	1000	4	avgt	10	0,025 ± 0,006		ms/op
MyBenchmark.mmReadIntFile	1000	1000	5	avgt	10	0,026 ± 0,007		ms/op
MyBenchmark.mmReadIntFile	1000	1000	10	avgt	10	0,037 ± 0,020		ms/op
MyBenchmark.mmReadIntFile	1000	1000	15	avgt	10	0,049 ± 0,024		ms/op
MyBenchmark.mmReadIntFile	1000	1000	20	avgt	10	0,065 ± 0,045		ms/op
MyBenchmark.mmReadIntFile	1000	1000	25	avgt	10	0,061 ± 0,046		ms/op
MyBenchmark.mmReadIntFile	1000	1000	30	avgt	10	0,060 ± 0,034		ms/op
MyBenchmark.readIntFile	N/A	1000	1	avgt	10	3,439 ± 0,127		ms/op
MyBenchmark.readIntFile	N/A	1000	2	avgt	10	3,530 ± 0,305		ms/op
MyBenchmark.readIntFile	N/A	1000	3	avgt	10	3,440 ± 0,208		ms/op
MyBenchmark.readIntFile	N/A	1000	4	avgt	10	3,421 ± 0,062		ms/op
MyBenchmark.readIntFile	N/A	1000	5	avgt	10	3,403 ± 0,087		ms/op
MyBenchmark.readIntFile	N/A	1000	10	avgt	10	3,492 ± 0,149		ms/op
MyBenchmark.readIntFile	N/A	1000	15	avgt	10	3,512 ± 0,141		ms/op
MyBenchmark.readIntFile	N/A	1000	20	avgt	10	3,531 ± 0,172		ms/op
MyBenchmark.readIntFile	N/A	1000	25	avgt	10	3,684 ± 0,393		ms/op
MyBenchmark.readIntFile	N/A	1000	30	avgt	10	3,515 ± 0,149		ms/op

Figure 10: Test N = 1.000

The first run highlights the fact

For the second run, the data volume is set to 10.000 32-bits integers.

```

Result "com.ulb.psk.MyBenchmark.readIntFile":
  35,053 ±(99.9%) 1,465 ms/op [Average]
  (min, avg, max) = (33,596, 35,053, 36,401), stdev = 0,969
  CI (99.9%): [33,588, 36,518] (assumes normal distribution)

# Run complete. Total time: 00:07:12

Benchmark      (bufferSize) (dataVolume) (filesToRead) Mode Cnt Score Error Units
MyBenchmark.mmReadIntFile 1000 10000 1  avgt 10 0,051 ± 0,005 ms/op
MyBenchmark.mmReadIntFile 1000 10000 2  avgt 10 0,064 ± 0,012 ms/op
MyBenchmark.mmReadIntFile 1000 10000 3  avgt 10 0,078 ± 0,007 ms/op
MyBenchmark.mmReadIntFile 1000 10000 4  avgt 10 0,097 ± 0,010 ms/op
MyBenchmark.mmReadIntFile 1000 10000 5  avgt 10 0,116 ± 0,012 ms/op
MyBenchmark.mmReadIntFile 1000 10000 10 avgt 10 0,178 ± 0,003 ms/op
MyBenchmark.mmReadIntFile 1000 10000 15 avgt 10 0,161 ± 0,012 ms/op
MyBenchmark.mmReadIntFile 1000 10000 20 avgt 10 0,175 ± 0,033 ms/op
MyBenchmark.mmReadIntFile 1000 10000 25 avgt 10 0,189 ± 0,054 ms/op
MyBenchmark.mmReadIntFile 1000 10000 30 avgt 10 0,208 ± 0,067 ms/op
MyBenchmark.readIntFile N/A 10000 1  avgt 10 3,521 ± 0,072 ms/op
MyBenchmark.readIntFile N/A 10000 2  avgt 10 7,009 ± 0,179 ms/op
MyBenchmark.readIntFile N/A 10000 3  avgt 10 10,444 ± 0,216 ms/op
MyBenchmark.readIntFile N/A 10000 4  avgt 10 13,907 ± 0,186 ms/op
MyBenchmark.readIntFile N/A 10000 5  avgt 10 17,129 ± 0,514 ms/op
MyBenchmark.readIntFile N/A 10000 10 avgt 10 34,641 ± 0,962 ms/op
MyBenchmark.readIntFile N/A 10000 15 avgt 10 34,300 ± 0,832 ms/op
MyBenchmark.readIntFile N/A 10000 20 avgt 10 34,830 ± 1,022 ms/op
MyBenchmark.readIntFile N/A 10000 25 avgt 10 34,581 ± 0,768 ms/op
MyBenchmark.readIntFile N/A 10000 30 avgt 10 35,053 ± 1,465 ms/op

```

Figure 11: Test N = 10.000

For the third run, the data volume is set to 100.000 32-bits integers.

```

Result "com.ulb.psk.MyBenchmark.readIntFile":
  104,734 ±(99.9%) 1,479 ms/op [Average]
  (min, avg, max) = (103,371, 104,734, 106,095), stdev = 0,979
  CI (99.9%): [103,254, 106,213] (assumes normal distribution)

# Run complete. Total time: 00:07:06

Benchmark      (bufferSize) (dataVolume) (filesToRead) Mode Cnt Score Error Units
MyBenchmark.mmReadIntFile 1000 100000 1  avgt 10 0,237 ± 0,079 ms/op
MyBenchmark.mmReadIntFile 1000 100000 2  avgt 10 0,285 ± 0,016 ms/op
MyBenchmark.mmReadIntFile 1000 100000 3  avgt 10 0,344 ± 0,077 ms/op
MyBenchmark.mmReadIntFile 1000 100000 4  avgt 10 0,350 ± 0,039 ms/op
MyBenchmark.mmReadIntFile 1000 100000 5  avgt 10 0,347 ± 0,030 ms/op
MyBenchmark.mmReadIntFile 1000 100000 10 avgt 10 0,406 ± 0,022 ms/op
MyBenchmark.mmReadIntFile 1000 100000 15 avgt 10 0,495 ± 0,053 ms/op
MyBenchmark.mmReadIntFile 1000 100000 20 avgt 10 0,545 ± 0,048 ms/op
MyBenchmark.mmReadIntFile 1000 100000 25 avgt 10 0,643 ± 0,069 ms/op
MyBenchmark.mmReadIntFile 1000 100000 30 avgt 10 0,719 ± 0,065 ms/op
MyBenchmark.readIntFile N/A 100000 1  avgt 10 3,673 ± 0,092 ms/op
MyBenchmark.readIntFile N/A 100000 2  avgt 10 7,254 ± 0,200 ms/op
MyBenchmark.readIntFile N/A 100000 3  avgt 10 10,811 ± 0,343 ms/op
MyBenchmark.readIntFile N/A 100000 4  avgt 10 14,240 ± 0,396 ms/op
MyBenchmark.readIntFile N/A 100000 5  avgt 10 17,713 ± 0,435 ms/op
MyBenchmark.readIntFile N/A 100000 10 avgt 10 34,824 ± 0,988 ms/op
MyBenchmark.readIntFile N/A 100000 15 avgt 10 52,464 ± 1,598 ms/op
MyBenchmark.readIntFile N/A 100000 20 avgt 10 70,205 ± 2,520 ms/op
MyBenchmark.readIntFile N/A 100000 25 avgt 10 90,564 ± 3,898 ms/op
MyBenchmark.readIntFile N/A 100000 30 avgt 10 104,734 ± 1,479 ms/op

```

Figure 12: Test N = 100.000

For the fourth run, the data volume is set to 1.000.000 32-bits integers.

```
Result "com.ulb.psk.MyBenchmark.readIntFile":
107,057 ±(99.9%) 2,537 ms/op [Average]
(min, avg, max) = (104,926, 107,057, 109,801), stdev = 1,678
CI (99.9%): [104,520, 109,594] (assumes normal distribution)
```

```
# Run complete. Total time: 00:07:12
```

Benchmark	(bufferSize)	(dataVolume)	(filesToRead)	Mode	Cnt	Score	Error	Units
MyBenchmark.mmReadIntFile	1000	1000000	1	avgt	10	2,225 ± 0,444		ms/op
MyBenchmark.mmReadIntFile	1000	1000000	2	avgt	10	2,688 ± 0,228		ms/op
MyBenchmark.mmReadIntFile	1000	1000000	3	avgt	10	2,680 ± 0,234		ms/op
MyBenchmark.mmReadIntFile	1000	1000000	4	avgt	10	2,777 ± 0,212		ms/op
MyBenchmark.mmReadIntFile	1000	1000000	5	avgt	10	2,802 ± 0,466		ms/op
MyBenchmark.mmReadIntFile	1000	1000000	10	avgt	10	2,549 ± 0,053		ms/op
MyBenchmark.mmReadIntFile	1000	1000000	15	avgt	10	2,713 ± 0,174		ms/op
MyBenchmark.mmReadIntFile	1000	1000000	20	avgt	10	2,873 ± 0,199		ms/op
MyBenchmark.mmReadIntFile	1000	1000000	25	avgt	10	2,753 ± 0,092		ms/op
MyBenchmark.mmReadIntFile	1000	1000000	30	avgt	10	2,850 ± 0,061		ms/op
MyBenchmark.readIntFile	N/A	1000000	1	avgt	10	5,449 ± 0,120		ms/op
MyBenchmark.readIntFile	N/A	1000000	2	avgt	10	9,314 ± 0,256		ms/op
MyBenchmark.readIntFile	N/A	1000000	3	avgt	10	12,842 ± 0,393		ms/op
MyBenchmark.readIntFile	N/A	1000000	4	avgt	10	16,139 ± 0,234		ms/op
MyBenchmark.readIntFile	N/A	1000000	5	avgt	10	19,549 ± 0,573		ms/op
MyBenchmark.readIntFile	N/A	1000000	10	avgt	10	37,462 ± 1,151		ms/op
MyBenchmark.readIntFile	N/A	1000000	15	avgt	10	54,486 ± 1,582		ms/op
MyBenchmark.readIntFile	N/A	1000000	20	avgt	10	71,857 ± 1,964		ms/op
MyBenchmark.readIntFile	N/A	1000000	25	avgt	10	89,376 ± 1,911		ms/op
MyBenchmark.readIntFile	N/A	1000000	30	avgt	10	107,057 ± 2,537		ms/op

Figure 13: Test N = 1.000.000

For the fifth run, the data volume is set to 5.000.000 32-bits integers.

```

Result "com.ulb.psk.MyBenchmark.readIntFile":
136,959 ±(99.9%) 2,997 ms/op [Average]
(min, avg, max) = (132,980, 136,959, 138,862), stdev = 1,983
CI (99.9%): [133,962, 139,956] (assumes normal distribution)

```

```
# Run complete. Total time: 00:07:18
```

Benchmark	(bufferSize)	(dataVolume)	(filesToRead)	Mode	Cnt	Score	Error	Units
MyBenchmark.mmReadIntFile	1000	5000000	1	avgt	10	11,745 ± 0,874		ms/op
MyBenchmark.mmReadIntFile	1000	5000000	2	avgt	10	14,774 ± 1,135		ms/op
MyBenchmark.mmReadIntFile	1000	5000000	3	avgt	10	14,993 ± 1,169		ms/op
MyBenchmark.mmReadIntFile	1000	5000000	4	avgt	10	14,461 ± 0,863		ms/op
MyBenchmark.mmReadIntFile	1000	5000000	5	avgt	10	14,080 ± 0,507		ms/op
MyBenchmark.mmReadIntFile	1000	5000000	10	avgt	10	13,530 ± 0,579		ms/op
MyBenchmark.mmReadIntFile	1000	5000000	15	avgt	10	13,593 ± 0,484		ms/op
MyBenchmark.mmReadIntFile	1000	5000000	20	avgt	10	13,700 ± 0,497		ms/op
MyBenchmark.mmReadIntFile	1000	5000000	25	avgt	10	13,577 ± 0,252		ms/op
MyBenchmark.mmReadIntFile	1000	5000000	30	avgt	10	13,880 ± 0,350		ms/op
MyBenchmark.readIntFile	N/A	5000000	1	avgt	10	15,141 ± 0,280		ms/op
MyBenchmark.readIntFile	N/A	5000000	2	avgt	10	22,270 ± 0,537		ms/op
MyBenchmark.readIntFile	N/A	5000000	3	avgt	10	26,095 ± 0,716		ms/op
MyBenchmark.readIntFile	N/A	5000000	4	avgt	10	29,826 ± 0,280		ms/op
MyBenchmark.readIntFile	N/A	5000000	5	avgt	10	34,321 ± 1,144		ms/op
MyBenchmark.readIntFile	N/A	5000000	10	avgt	10	55,156 ± 2,156		ms/op
MyBenchmark.readIntFile	N/A	5000000	15	avgt	10	76,161 ± 1,877		ms/op
MyBenchmark.readIntFile	N/A	5000000	20	avgt	10	96,675 ± 1,870		ms/op
MyBenchmark.readIntFile	N/A	5000000	25	avgt	10	116,978 ± 2,687		ms/op
MyBenchmark.readIntFile	N/A	5000000	30	avgt	10	136,959 ± 2,997		ms/op

Figure 14: Test N = 5.000.000

Now that different value of N and k have been tried there is just the size of the buffer to play with. In order to test the impact of the buffer size, B is set to multiple value; 10.000, 50.000, 100.000, 1.000.000 and 10.000.000. More over in order to test this with a big size in input, N is set to 10.000.000 32-bits integers. This lead to the following results :

```

Result "com.ulb.psk.MyBenchmark.readIntFile":
18719,727 ±(99.9%) 297,961 ms/op [Average]
(min, avg, max) = (18304,200, 18719,727, 18987,037), stdev = 197,083
CI (99.9%): [18421,766, 19017,689] (assumes normal distribution)

```

```
# Run complete. Total time: 01:32:01
```

Benchmark	(bufferSize)	(dataVolume)	(filesToRead)	Mode	Cnt	Score	Error	Units
MyBenchmark.mmReadIntFile	10000	5000000	1	avgt	10	12,875 ±	2,442	ms/op
MyBenchmark.mmReadIntFile	10000	5000000	2	avgt	10	15,441 ±	1,542	ms/op
MyBenchmark.mmReadIntFile	10000	5000000	3	avgt	10	32,350 ±	26,112	ms/op
MyBenchmark.mmReadIntFile	10000	5000000	4	avgt	10	17,152 ±	2,315	ms/op
MyBenchmark.mmReadIntFile	10000	5000000	5	avgt	10	19,991 ±	6,075	ms/op
MyBenchmark.mmReadIntFile	10000	5000000	10	avgt	10	18,611 ±	5,316	ms/op
MyBenchmark.mmReadIntFile	10000	5000000	15	avgt	10	16,528 ±	1,317	ms/op
MyBenchmark.mmReadIntFile	10000	5000000	20	avgt	10	17,380 ±	0,564	ms/op
MyBenchmark.mmReadIntFile	10000	5000000	25	avgt	10	18,095 ±	0,730	ms/op
MyBenchmark.mmReadIntFile	10000	5000000	30	avgt	10	18,845 ±	0,543	ms/op
MyBenchmark.mmReadIntFile	50000	5000000	1	avgt	10	12,279 ±	0,660	ms/op
MyBenchmark.mmReadIntFile	50000	5000000	2	avgt	10	15,800 ±	0,218	ms/op
MyBenchmark.mmReadIntFile	50000	5000000	3	avgt	10	16,794 ±	0,418	ms/op
MyBenchmark.mmReadIntFile	50000	5000000	4	avgt	10	17,507 ±	0,409	ms/op
MyBenchmark.mmReadIntFile	50000	5000000	5	avgt	10	18,651 ±	0,453	ms/op
MyBenchmark.mmReadIntFile	50000	5000000	10	avgt	10	22,391 ±	0,704	ms/op
MyBenchmark.mmReadIntFile	50000	5000000	15	avgt	10	26,194 ±	1,077	ms/op
MyBenchmark.mmReadIntFile	50000	5000000	20	avgt	10	31,405 ±	0,746	ms/op
MyBenchmark.mmReadIntFile	50000	5000000	25	avgt	10	37,415 ±	3,607	ms/op
MyBenchmark.mmReadIntFile	50000	5000000	30	avgt	10	40,810 ±	2,984	ms/op
MyBenchmark.mmReadIntFile	100000	5000000	1	avgt	10	13,366 ±	0,812	ms/op
MyBenchmark.mmReadIntFile	100000	5000000	2	avgt	10	17,991 ±	1,183	ms/op
MyBenchmark.mmReadIntFile	100000	5000000	3	avgt	10	19,205 ±	0,494	ms/op
MyBenchmark.mmReadIntFile	100000	5000000	4	avgt	10	21,072 ±	0,637	ms/op
MyBenchmark.mmReadIntFile	100000	5000000	5	avgt	10	23,013 ±	0,735	ms/op
MyBenchmark.mmReadIntFile	100000	5000000	10	avgt	10	31,194 ±	1,199	ms/op
MyBenchmark.mmReadIntFile	100000	5000000	15	avgt	10	39,261 ±	0,968	ms/op
MyBenchmark.mmReadIntFile	100000	5000000	20	avgt	10	49,326 ±	1,684	ms/op
MyBenchmark.mmReadIntFile	100000	5000000	25	avgt	10	62,251 ±	3,664	ms/op
MyBenchmark.mmReadIntFile	100000	5000000	30	avgt	10	77,889 ±	11,909	ms/op
MyBenchmark.mmReadIntFile	1000000	5000000	1	avgt	10	31,471 ±	1,911	ms/op
MyBenchmark.mmReadIntFile	1000000	5000000	2	avgt	10	54,484 ±	2,044	ms/op
MyBenchmark.mmReadIntFile	1000000	5000000	3	avgt	10	74,664 ±	4,196	ms/op
MyBenchmark.mmReadIntFile	1000000	5000000	4	avgt	10	98,325 ±	7,811	ms/op
MyBenchmark.mmReadIntFile	1000000	5000000	5	avgt	10	118,656 ±	9,817	ms/op
MyBenchmark.mmReadIntFile	1000000	5000000	10	avgt	10	110,955 ±	10,422	ms/op
MyBenchmark.mmReadIntFile	1000000	5000000	15	avgt	10	111,139 ±	24,063	ms/op
MyBenchmark.mmReadIntFile	1000000	5000000	20	avgt	10	118,878 ±	21,220	ms/op
MyBenchmark.mmReadIntFile	1000000	5000000	25	avgt	10	108,990 ±	10,534	ms/op
MyBenchmark.mmReadIntFile	1000000	5000000	30	avgt	10	124,856 ±	23,112	ms/op
MyBenchmark.mmReadIntFile	10000000	5000000	1	avgt	10	103,828 ±	16,509	ms/op
MyBenchmark.mmReadIntFile	10000000	5000000	2	avgt	10	119,951 ±	14,541	ms/op
MyBenchmark.mmReadIntFile	10000000	5000000	3	avgt	10	139,568 ±	32,350	ms/op
MyBenchmark.mmReadIntFile	10000000	5000000	4	avgt	10	124,137 ±	23,244	ms/op
MyBenchmark.mmReadIntFile	10000000	5000000	5	avgt	10	127,802 ±	13,124	ms/op
MyBenchmark.mmReadIntFile	10000000	5000000	10	avgt	10	127,057 ±	8,898	ms/op
MyBenchmark.mmReadIntFile	10000000	5000000	15	avgt	10	107,448 ±	24,572	ms/op
MyBenchmark.mmReadIntFile	10000000	5000000	20	avgt	10	124,500 ±	18,767	ms/op
MyBenchmark.mmReadIntFile	10000000	5000000	25	avgt	10	160,425 ±	58,940	ms/op
MyBenchmark.mmReadIntFile	10000000	5000000	30	avgt	10	169,808 ±	74,346	ms/op
MyBenchmark.readIntFile	N/A	5000000	1	avgt	10	22195,341 ±	2662,440	ms/op
MyBenchmark.readIntFile	N/A	5000000	2	avgt	10	20444,796 ±	83,715	ms/op
MyBenchmark.readIntFile	N/A	5000000	3	avgt	10	24370,901 ±	4289,356	ms/op
MyBenchmark.readIntFile	N/A	5000000	4	avgt	10	24041,832 ±	1807,636	ms/op
MyBenchmark.readIntFile	N/A	5000000	5	avgt	10	23285,622 ±	1473,546	ms/op
MyBenchmark.readIntFile	N/A	5000000	10	avgt	10	23823,957 ±	4217,155	ms/op
MyBenchmark.readIntFile	N/A	5000000	15	avgt	10	20903,211 ±	923,341	ms/op
MyBenchmark.readIntFile	N/A	5000000	20	avgt	10	21964,376 ±	1567,025	ms/op
MyBenchmark.readIntFile	N/A	5000000	25	avgt	10	20506,977 ±	811,356	ms/op
MyBenchmark.readIntFile	N/A	5000000	30	avgt	10	18719,727 ±	297,961	ms/op

Figure 15: Test on buffer size

2.7 Discussion of expected behavior vs experimental observations

Those benchmarks seem to give results that are the same as the expected behavior. In fact the First implementation is from far away the one that take the more time to read and write a lot of data. The second and the third one have almost the same time of execution and the fourth one is almost all the time getting better result. Thus the stream implementation to use for the external multi-way merge-sort algorithm is the fourth mechanism. So the algorithm will use memory mapping mechanism to read and write faster.

3 Observations on multi-way merge sort

3.1 Implementation

First a multi way merging algorithm needs to be implemented. The algorithm consists of a d-way merging algorithm that, given d sorted input streams of 32-bit integer, creates a single output stream containing the elements from the input stream in sorted order.

In order to give to the algorithm the d sorted input streams, an ArrayList is filled up with ArrayList that represent each one an input stream. The merging is using a priority queue in order to obtain the next element to be output at all times.

Then once the d-way merging algorithm is done, a new implementation needs to be done. In fact the new algorithm to implement is an external memory multi-way merge-sort algorithm for sorting 32-bit integers. The algorithm takes as input an input file and the following parameters:

- M - the size of the main memory available during the first sort phase, in number of 32-bit integers;
- d - the number of streams to merge in one pass in the later sort phases

Moreover N is defined as being the size of the input file, measured in number of 32-bit integers. The program should be able to sort in a particular way.

- The program needs to read the input file and be able to split in into N/M streams. In our implementation, the program creates first N/M files then he fills each of them with a sorted stream. In order to sort all the stream the method sort from java.util.Collections is used. According to the documentation this sort algorithm offers guaranteed $\log(n)$ performance.
- In order to store the references to the N/M streams a queue is used.
- At this point the program needs to merge the d first streams in the queue and put the resulting stream at the end of the queue. This has to be

repeated until the number of streams that remain in the queue is under d .

TCB:

In this project two-pass algorithms are used, where data from the operand relations is read into main memory, processed, written out to disk again, and then reread from disk to complete the operation. Two passes are usually enough, even for very large relations and generalizing it to more than two passes is not hard. It begins with an implementation of the sorting operator r it divides the relation R for which $B(R)$ greater than M into chunks of size M , sort them, and then process the sorted sublists. It requires only one block of each sorted sublist in main memory.

Two-Phase Multiway Merge-Sort (TPMMS) sorts very large relations in two passes. M is used, as the main memory buffers to sort. TPMMS fills repeatedly the M buffers with new tuples from R (Relation) uses main-memory sorting algorithm and sorts them and then writes out each sorted sublist to secondary storage after that it merges at most $M - 1$ sorted sublists to secondary storage, which limits the size of R . One input block to each sorted sublist and one block are allocated to the output. The smallest key among the first remaining elements of all the lists with is found and this comparison is done in main memory. A linear search is used for taking a number of machine instructions proportional to the number of sublists. A method based on "priority queues"² is used that takes time proportional to the logarithm of the number of sublists to find the smallest element then the smallest element is moved to the first available position of the output block. If the output block is full, write it to disk and reinitialize the same buffer in main memory to hold the next output block. If the block from which the smallest element was just taken is now exhausted of records, read the next block from the same sorted sublist into the same buffer that was used for the block just exhausted. If no blocks remain, then leave its buffer empty and do not consider elements from that list in any further competition for smallest remaining elements.

In order for TPMMS to work, there must be no more than $M - 1$ sublists. Suppose R fits on B blocks. Since each sublist consists of M blocks, the number of sublists is $\lceil B/M \rceil$. thus it requires $\lceil B/M \rceil \leq M - 1$, or $B \leq M(M - 1)$ (or about $B \leq M^2$). The algorithm requires us to read B blocks in the first pass, and another B disk I/O's to write the sorted sublists. The sorted sublists are each read again in the second pass, resulting in a total of $3B$ disk I/O's. If, as is customary, The cost of writing the result to disk is not counted (since the result may be pipelined and never written to disk), then $2B$ is all that the sorting operator r requires. However, if storing the result on disk is needed to store the result on disk, then the requirement is $4B$.

2PMMS is sufficient to sort all but an incredibly large relation in two passes. However, if you have an even bigger relation, then the same idea can be applied recursively. Divide the relation into chunks of size $M(M - 1)$, use 2PMMS to sort each one, and then treat the resulting sorted lists as sublists for a third pass. The idea extends similarly to any number of passes.

3.2 Expected behavior

Before running any test, a cost formula needs to be build for this implementation. The cost estimates the total number of I/O's that need to be done in function of N , d and M .

- N - the size of the input file, measured in number of 32-bit integers
- d - the number of streams to merge in one pass in the later sort phases
- M - the size of the main memory available during the first sort phase, in number of 32-bit integers

According to the course of Prof. S.VANSUMMEREN, a demonstration of the cost formula is build. Here is the demonstration of the cost formula :

The input that need to be sort is a relation R with a size of N 32-bit integers. Let's introduce $N(R)$ as the number of blocks of the relations.

- In the first pass the program reads M blocks at the same time from the input relation R , sort these by means of the method sort from `java.util.Collections`, and write the sorted resulting sublist to disk. After the first pass there is $N(R)/M$ sorted sublists of M blocks each.
- In the following passes the program read d blocks from these sublists and merge them into larger sorted sublists. (After the second pass there is $N(R)/(M * d)$ sorted sublists of $M * d$ blocks each, after the third pass $N(R)/(M * d^2)$ sorted sublists, etc ..)
- The algorithm repeats until it obtains a single sorted sublist.

What is the complexity of this algorithms ?

- In each pass we read and write the entire input relation exactly once.
- There are $\lceil \log_M N(R) \rceil$ passes
- The total cost is hence $2N(R) \log_M N(R)$ I/O operations.

3.3 Experimental observations

In order to test the external multi way merge algorithms with the three parameters; N , M and d , a benchmark has been implemented.

The first experiment is to vary only N and to keep M and d fixed. During the run N is taking those following values : 1.000.000, 2.000.000, 4.000.000, 8.000.000.


```

Result "com.ulb.psk.benchmark.MyBenchmark.externalMerge":
90951,075 ±(99.9%) 66541,415 ms/op [Average]
(min, avg, max) = (63867,396, 90951,075, 210048,026), stdev = 44013,026
CI (99.9%): [24409,660, 157492,489] (assumes normal distribution)

# Run complete. Total time: 00:50:05

Benchmark      (d) (fileSize) (firstPhaseMemory) Mode Cnt      Score      Error Units
MyBenchmark.externalMerge 10  100000000      10000 avgt  10  7106,874 ± 610,026 ms/op
MyBenchmark.externalMerge 10  200000000      10000 avgt  10 15946,740 ± 859,903 ms/op
MyBenchmark.externalMerge 10  400000000      10000 avgt  10 39383,103 ± 2588,189 ms/op
MyBenchmark.externalMerge 10  800000000      10000 avgt  10 90951,075 ± 66541,415 ms/op

```

Figure 16: Test on N

The conclusion of this first experiment is quite obvious, because more big is the input file more it will take time to apply the algorithm.

The second experiment is to vary only M and to keep N and d fixed. During the run M is taking those following values : 10.000, 100.000, 1.000.000, 10.000.000.

```

Result "com.ulb.psk.benchmark.MyBenchmark.externalMerge":
31086,351 ±(99.9%) 2715,593 ms/op [Average]
(min, avg, max) = (28816,708, 31086,351, 33867,968), stdev = 1796,197
CI (99.9%): [28370,758, 33801,944] (assumes normal distribution)

# Run complete. Total time: 00:52:11

Benchmark      (d) (fileSize) (firstPhaseMemory) Mode Cnt      Score      Error Units
MyBenchmark.externalMerge 10  40000000      10000 avgt  10  38638,161 ± 2913,278 ms/op
MyBenchmark.externalMerge 10  40000000      100000 avgt  10 28104,357 ± 4896,591 ms/op
MyBenchmark.externalMerge 10  40000000      1000000 avgt  10 25602,758 ± 736,116 ms/op
MyBenchmark.externalMerge 10  40000000      10000000 avgt  10 31359,117 ± 1190,794 ms/op
MyBenchmark.externalMerge 10  40000000      100000000 avgt  10 31086,351 ± 2715,593 ms/op

```

Figure 17: Test on M

The result of the second experiment highlights that the best value for M is not the biggest one. In fact if M has a size of 100.000 it is faster than if has a size of 10.000.000. In this experiment the best size for M is 1.000.000.

The third experiment is to vary only d and to keep M and N fixed. During the run M is taking those following values : 2, 3, 4, 5, 10, 15, 20, 30, 40. The value of M is set regarding the last experiment, by taking the value that gave the best result. So M is set to 1.000.000.

```

Result "com.ulb.psk.benchmark.MyBenchmark.externalMerge":
  25763,539 ±(99.9%) 1953,576 ms/op [Average]
  (min, avg, max) = (24380,601, 25763,539, 27776,064), stdev = 1292,170
  CI (99.9%): [23809,963, 27717,116] (assumes normal distribution)

# Run complete. Total time: 01:29:02

Benchmark              (d)  (fileSize)  (firstPhaseMemory)  Mode  Cnt      Score      Error  Units
MyBenchmark.externalMerge  2    400000000      1000000      avgt   10  41562,686 ± 3558,276  ms/op
MyBenchmark.externalMerge  3    400000000      1000000      avgt   10  34105,524 ± 5112,707  ms/op
MyBenchmark.externalMerge  4    400000000      1000000      avgt   10  30946,335 ± 2335,822  ms/op
MyBenchmark.externalMerge  5    400000000      1000000      avgt   10  28549,205 ± 1511,995  ms/op
MyBenchmark.externalMerge 10    400000000      1000000      avgt   10  26583,125 ± 1431,166  ms/op
MyBenchmark.externalMerge 15    400000000      1000000      avgt   10  26878,014 ± 1022,072  ms/op
MyBenchmark.externalMerge 20    400000000      1000000      avgt   10  26447,133 ± 518,501  ms/op
MyBenchmark.externalMerge 30    400000000      1000000      avgt   10  26986,284 ± 601,210  ms/op
MyBenchmark.externalMerge 40    400000000      1000000      avgt   10  25763,539 ± 1953,576  ms/op

```

Figure 18: Test on d

The third experiment shows that more there is streams to merge in one pass in the later sort phases more faster is the algorithm to execute.

3.4 Discussion of expected behavior vs experimental observations

According to the expected behavior and the experimental observations the good choice for the parameters is to set a big buffer but not too big, 1.10^6 seems to be a good value. More over the number of streams to merge in one pass in the later phases must also be big order to reduce the time of execution.

4 Overall conclusion

During this project we learned about "Java NIO FileChannel", "JAVA NIO – Memory-Mapped File", "Testing the code performance with JMH tool", "Memory Mapped File in Java", "Revisiting File InputStream and Reader instantiation", "Priority Queue insert with priority" and "Implement K Way Merge Algorithm". We also learned to implement proper benchmarks, and so to use the JHM library.

List of Figures

1	Computer configuration	6
2	Free space on the computer	6
3	Test N = 1.000	7
4	Test N = 100.000	8
5	Test N = 1.000.000	9
6	Test N = 5.000.000	10
7	Test N = 7.000.000	11
8	1st Test on buffer size	13
9	2nd Test on buffer size	15
10	Test N = 1.000	16
11	Test N = 10.000	17
12	Test N = 100.000	17
13	Test N = 1.000.000	18
14	Test N = 5.000.000	19
15	Test on buffer size	20
16	Test on N	24
17	Test on M	24
18	Test on d	25

5 References

<http://tutorials.jenkov.com/java-nio/file-channel.html>
<https://examples.javacodegeeks.com/corejava/nio/filechannel/java-nio-channels-filechannel-example/>
<https://tutorials.techmytalk.com/2014/11/05/java-nio-memory-mapped-files/>
<https://www.codeproject.com/Tips/683614/Things-to-Know-about-Memory-Mapped-File-in-Java>
<https://www.developer.com/java/other/article.php/1548681/Introduction-to-Memory-Mapped-IO-in-Java.htm>
<https://www.javacodegeeks.com/2013/05/power-of-java-memorymapped-file.html>
<http://www.mathcs.emory.edu/~cheung/Courses/554/Syllabus/4query-exec/2-pass=TPMMS.html/>
<https://sortingalgorithms.quora.com/K-way-Merge-Algorithms>
<https://www.programcreek.com/2014/05/merge-k-sorted-arrays-in-java/>
<http://www.sanfoundry.com/java-program-k-way-merge-algorithm/>
<http://www.geeksforgeeks.org/merge-k-sorted-arrays/>
<http://www.mathcs.emory.edu/~cheung/Courses/554/Syllabus/4-query-exec/2-pass=TPMMS.html/>
https://en.wikipedia.org/wiki/K-Way_Merge_Algorithms/
https://www.javamex.com/tutorials/collections/priorityqueue_example_heapsort.shtml
<https://stackoverflow.com/questions/29380093/priorityqueue-insert-with-priority>
<http://tutorials.jenkov.com/java-nio/file-channel.html>
<https://tutorials.techmytalk.com/2014/11/05/java-nio-memory-mapped-files>
<https://www.javacodegeeks.com/2013/05/power-of-java-memorymapped-file.html>
<http://tutorials.jenkov.com/java-performance/jmh.html>
<https://arnaudroger.github.io/blog/2017/03/20/faster-reader-inpustream-in-java.html>
http://cs.ulb.ac.be/public/_media/teaching/infoh417/slides-lect6.pdf
<https://blog.goyello.com/2017/06/19/testing-code-performance-jmh-tool/>
<http://s-j.github.io/getting-started-with-jmh/>